

Docker Swarm

Blacksburg Docker Meetup

November 9, 2016



Quick review...

- What's an image?
- What's a container?
- What's Docker Compose?

Docker Compose Shortcomings

- Only runs containers on a single host
 - Works fantastic locally (development)
 - Single host = not preferred for prod

```
version: '2'
services:
  apache:
    image: php:7.0-apache
  redis:
    image: redis:3.0
  mysql:
    image: mysql:5.7
```




Quick V1 vs V2 Review...

- What's the difference??


```
version: '2'
services:
  apache:
    image: php:7.0-apache
  redis:
    image: redis:3.0
  mysql:
    image: mysql:5.7
```

```
apache:
  image: php:7.0-apache
redis:
  image: redis:3.0
mysql:
  image: mysql:5.7
```





Docker Compose Shortcomings

- Purely “fire and forget”




`docker-compose scale apache=3`





Enter the world of orchestration...




Introducing Docker Swarm

- Docker Swarm was pulled into the Docker Engine with Docker 1.12
 - There was a previous version of Swarm, but was a lot harder to use
- Completely optional to use
- BUT... comes with a ton of cool features
 - Swarm mode
 - Secure by default
 - Routing mesh
 - Service API

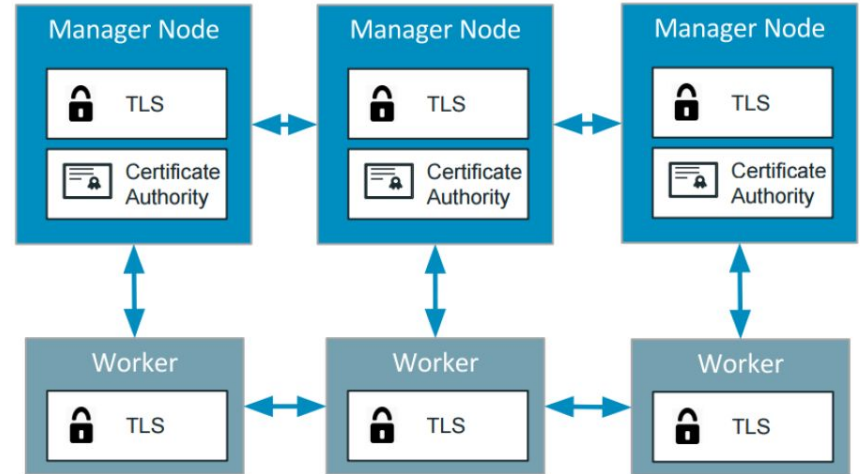
Swarm Mode

- Self-organizing and self-healing
- Completely decentralized architecture
 - No external data store required
 - No single points of failure



Secure by Default

- Cryptographic node identity
- Automatic encryption and mutual auth (TLS)
 - End-to-end TLS
- Automatic cert rotation
- External CA integration
- Built-in government-grade PKI



Routing Mesh

- Swarm-wide overlay network
- Container-native load-balancing
- DNS-based service discovery
- Kernel-only data path using IPVS
- Every node knows how to route to any other service in the swarm

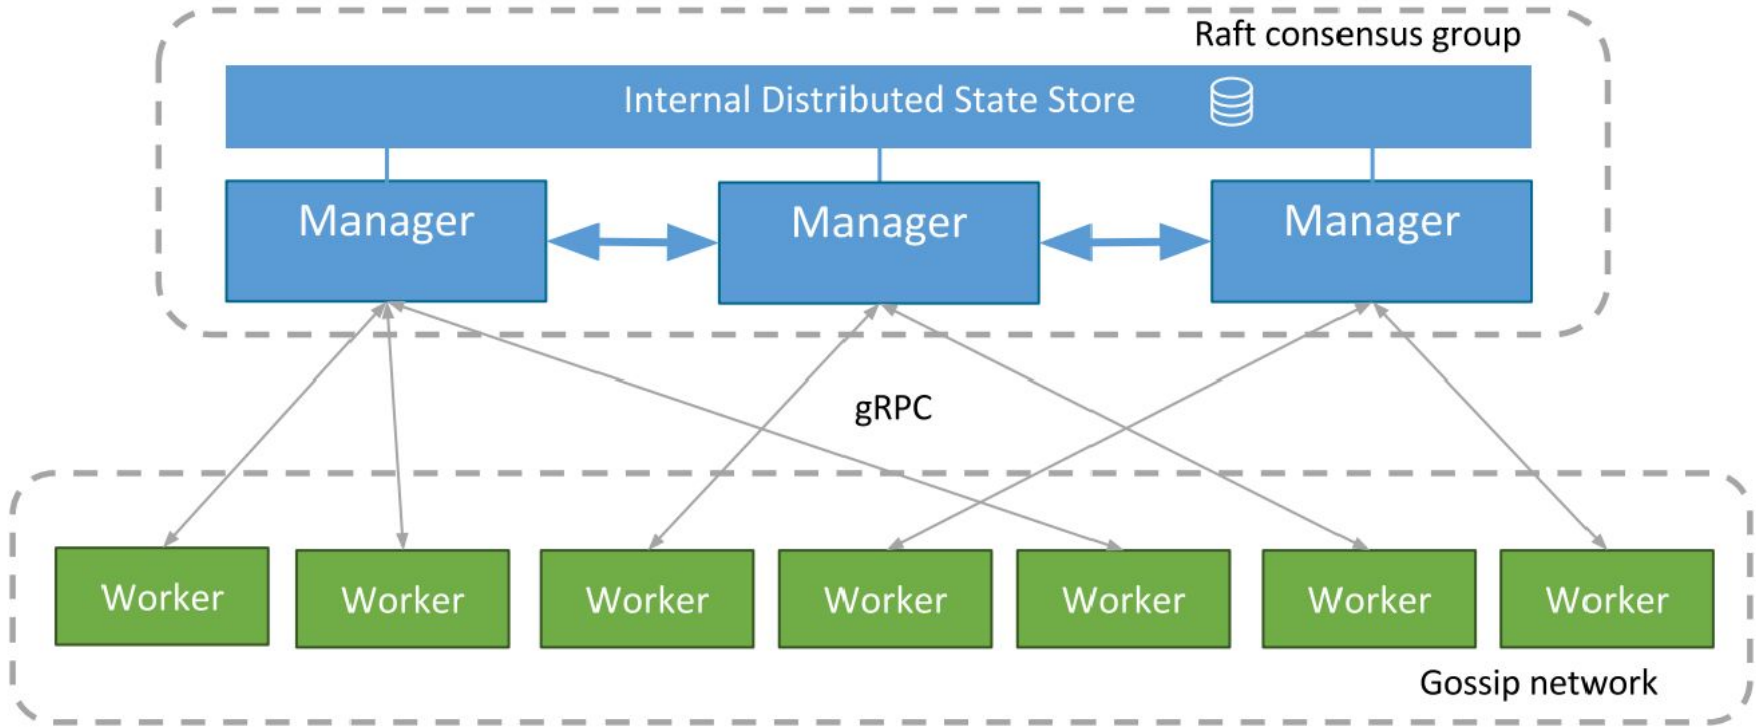


Docker Service API

- New CLI commands to set desired state
- Allows for...
 - Scaling
 - Rolling updates
 - Advanced scheduling
 - App-specific healthchecks
 - Rescheduling on node failure

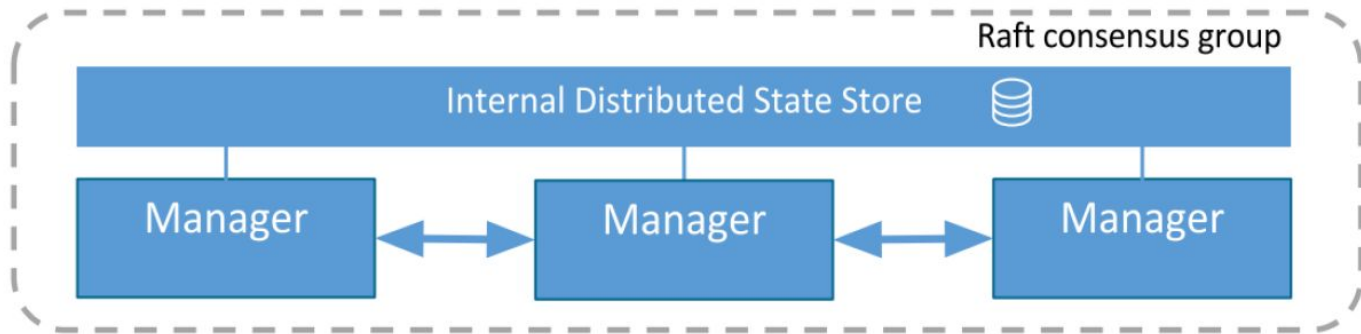


Internal Communications



Quorum Layer (Managers)

- Strongly consistent - holds desired state
- Simple to operate using Docker CLI
- Blazing fast (in-memory reads, domain specific indexing)
- Secure communications

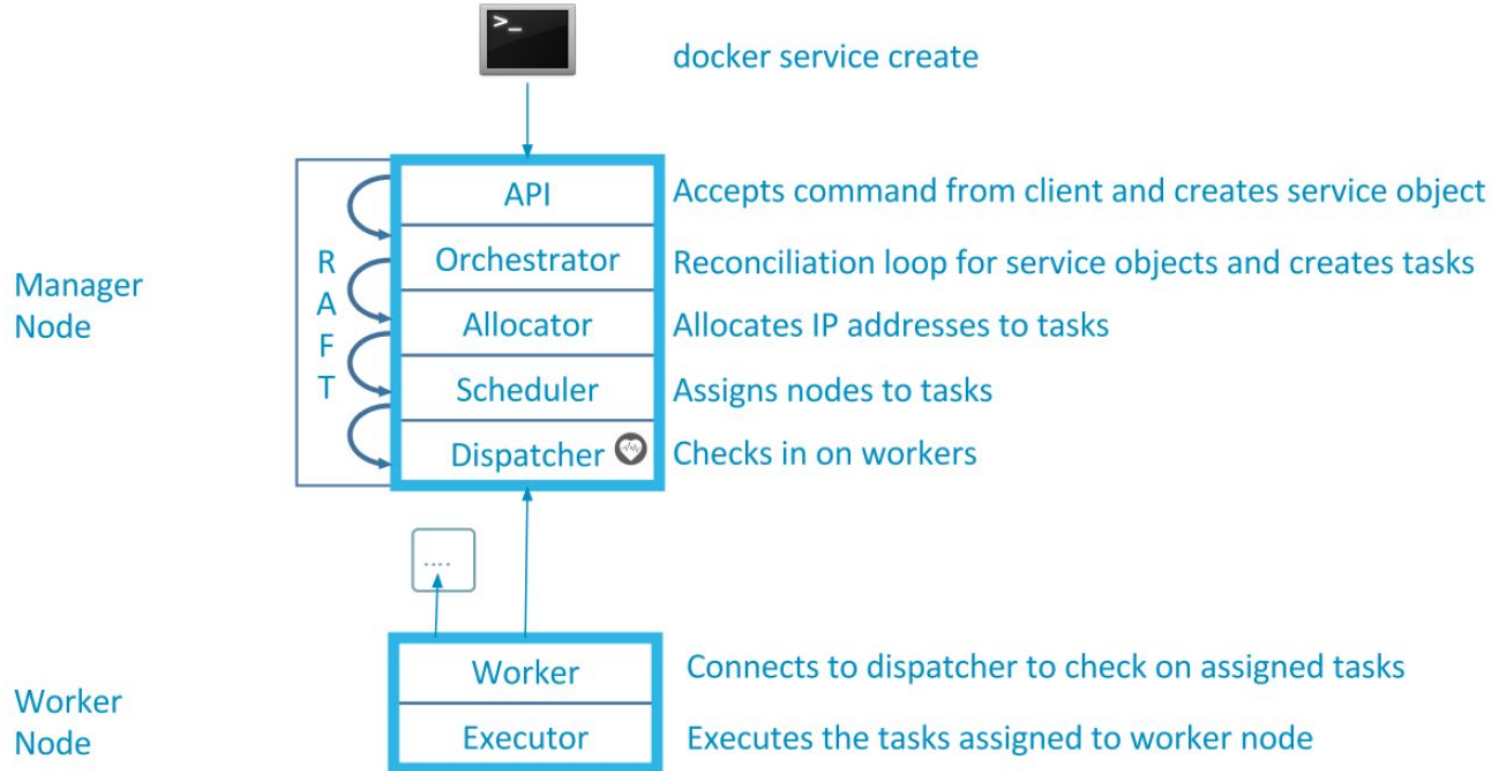


Work-to-Worker Gossip

- Eventually consistent: Routing mesh, load balancing rules
- High volume, P2P network
- Secure - symmetric encryption with key rotation



Under the Hood...



Let's see it in action...

Swarm initialization



Node-1


```
docker swarm init
```

Adding to the cluster




```
docker swarm join <IP of node-1>:2377
```

Adding to the cluster




```
docker swarm join <IP of node-1>:2377
```

Starting a service




```
docker service create --name app --network app -p 80:80 myapp
```

Another service



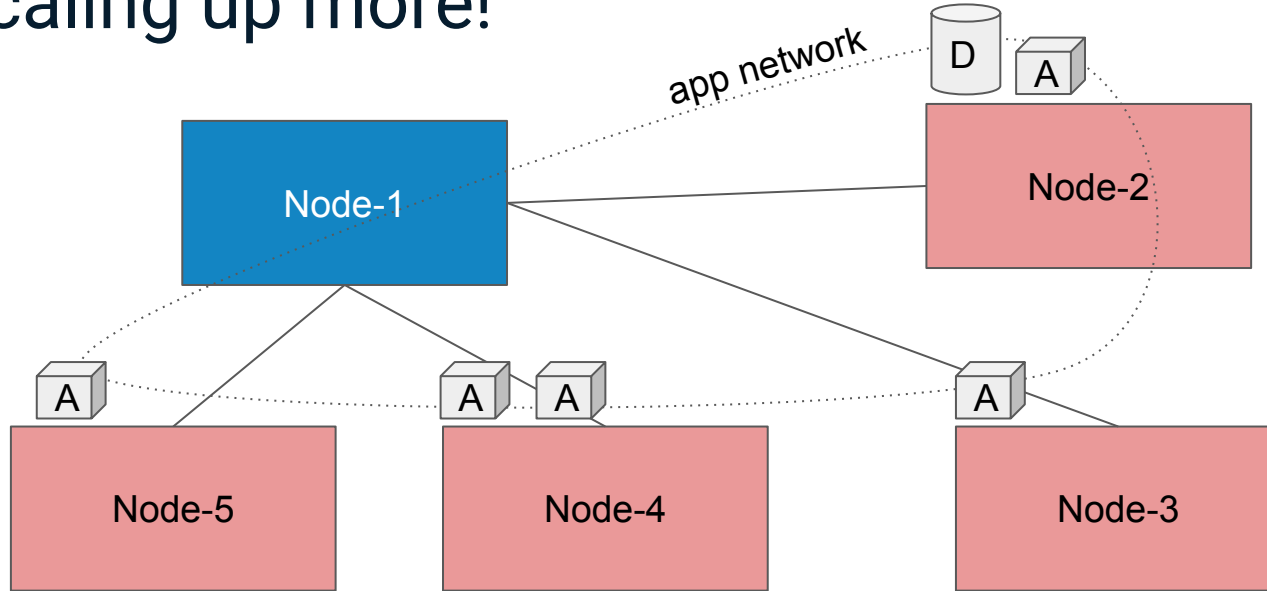
```
docker service create --name db --network app mysql
```

Scaling up



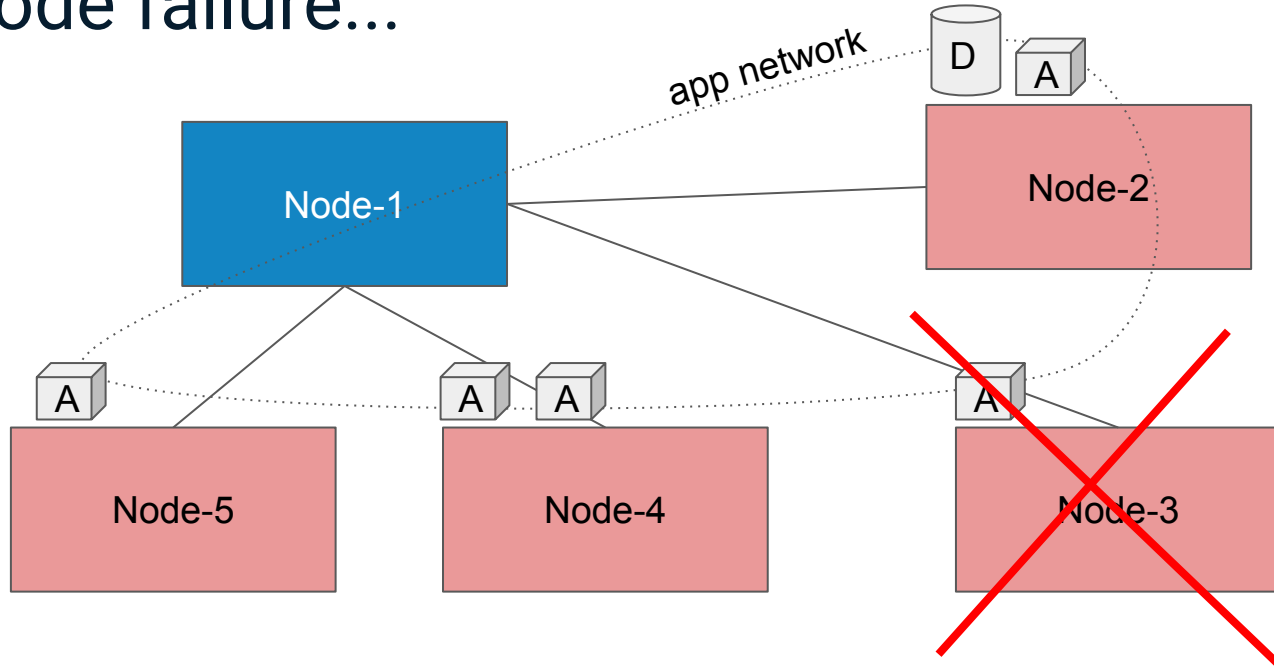
`docker service scale app=3`

Scaling up more!

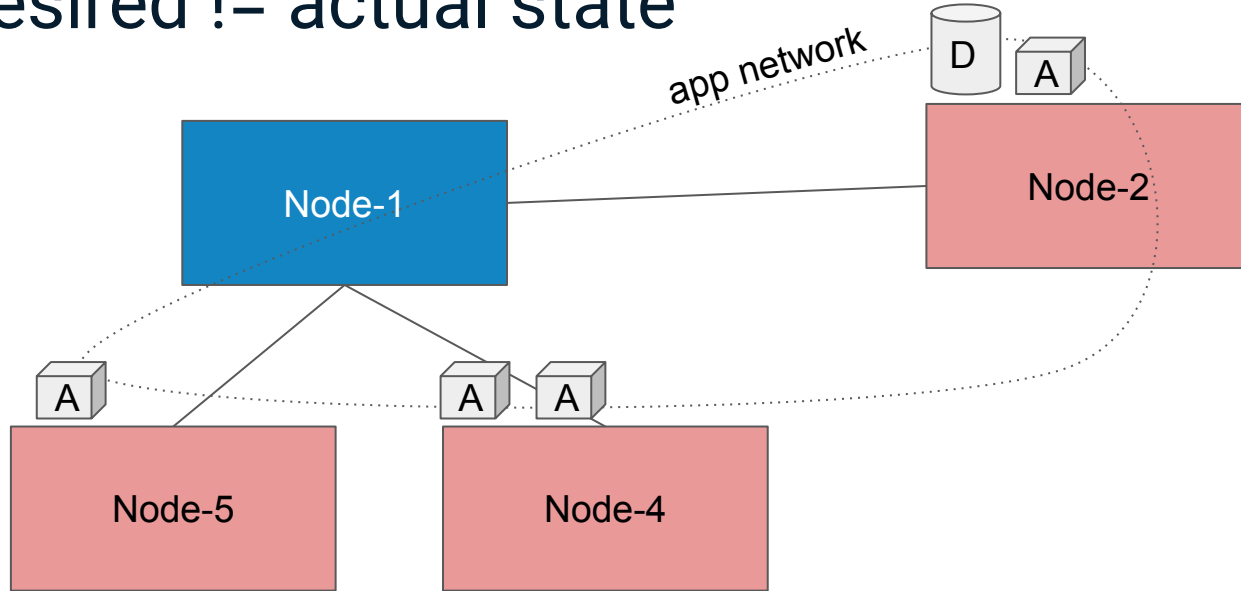


`docker service scale app=5`

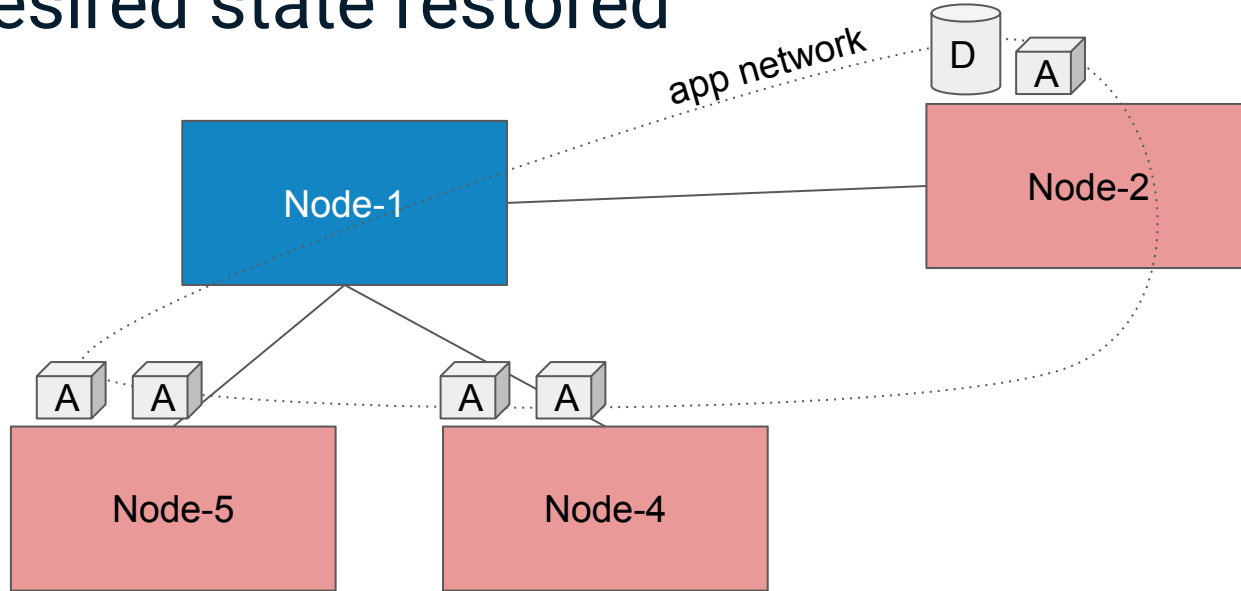
Node failure...



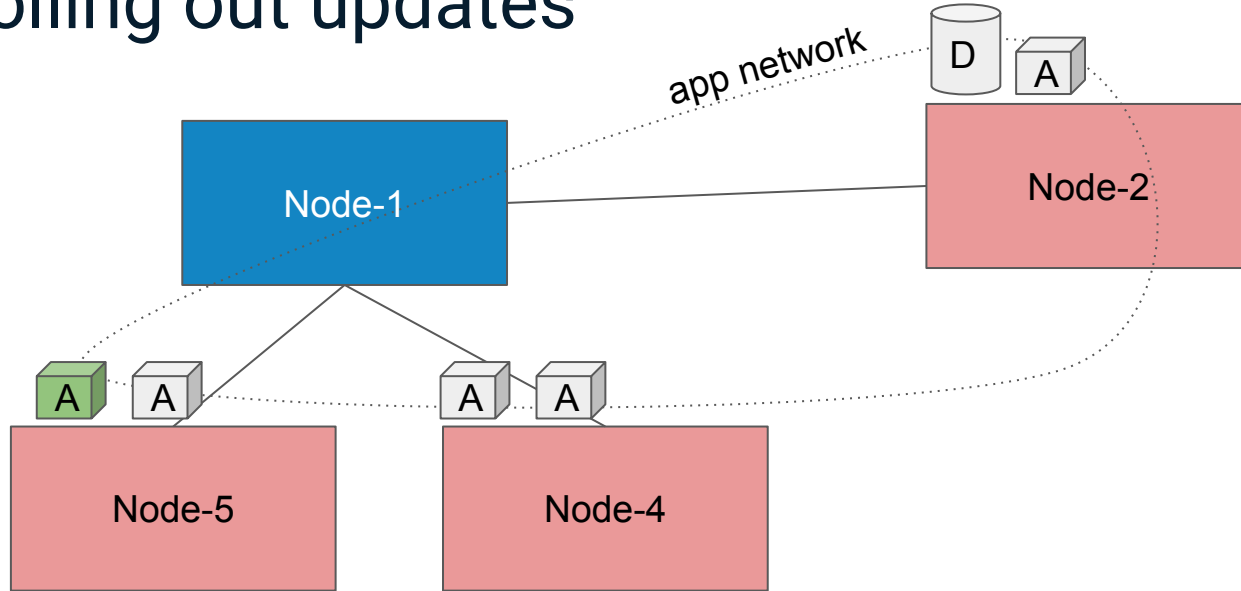
Desired != actual state



Desired state restored

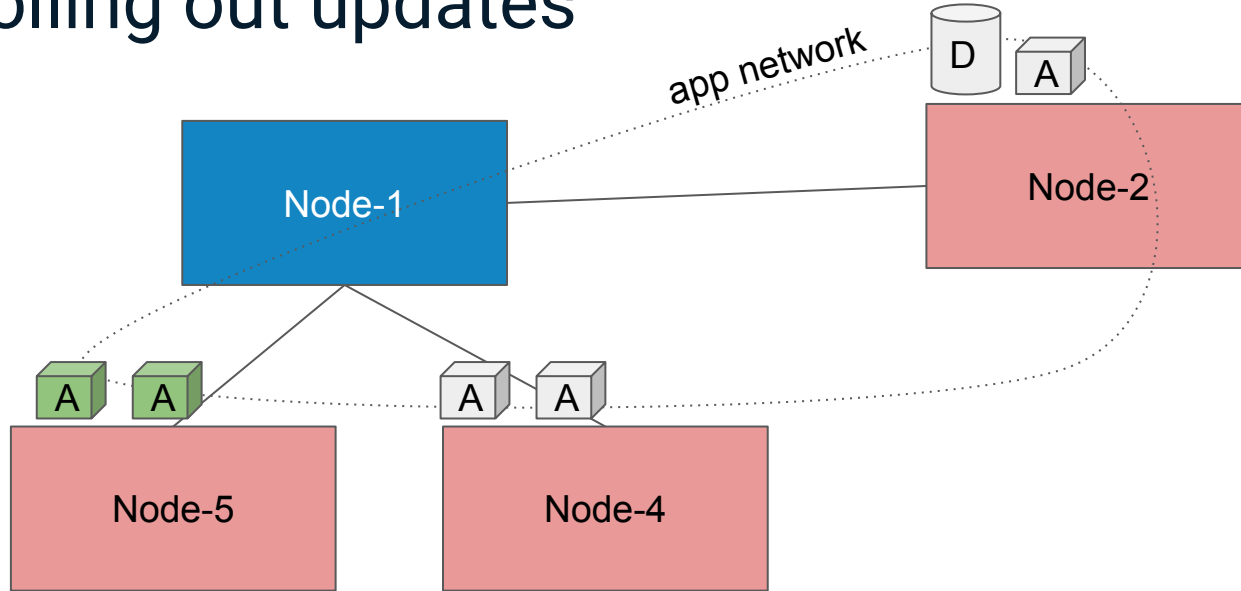


Rolling out updates



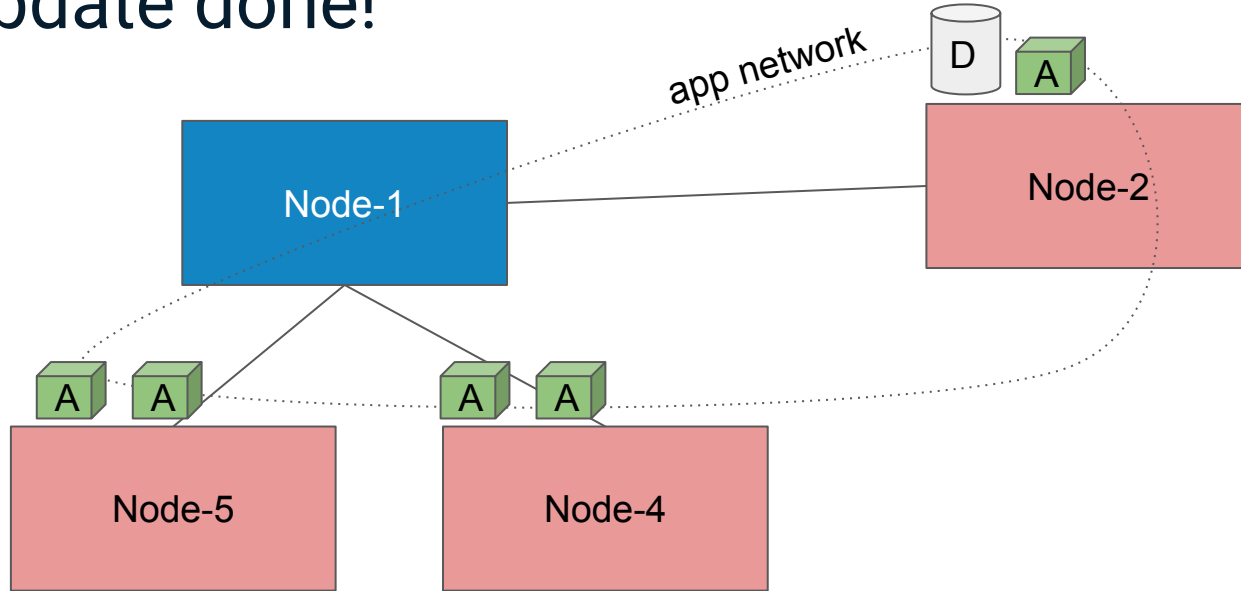
```
docker service update --image=app:1.1 --update-delay 10s app
```

Rolling out updates




Updates one, waits until it's in RUNNING state, then waits 10 seconds to move on to next task

Update done!

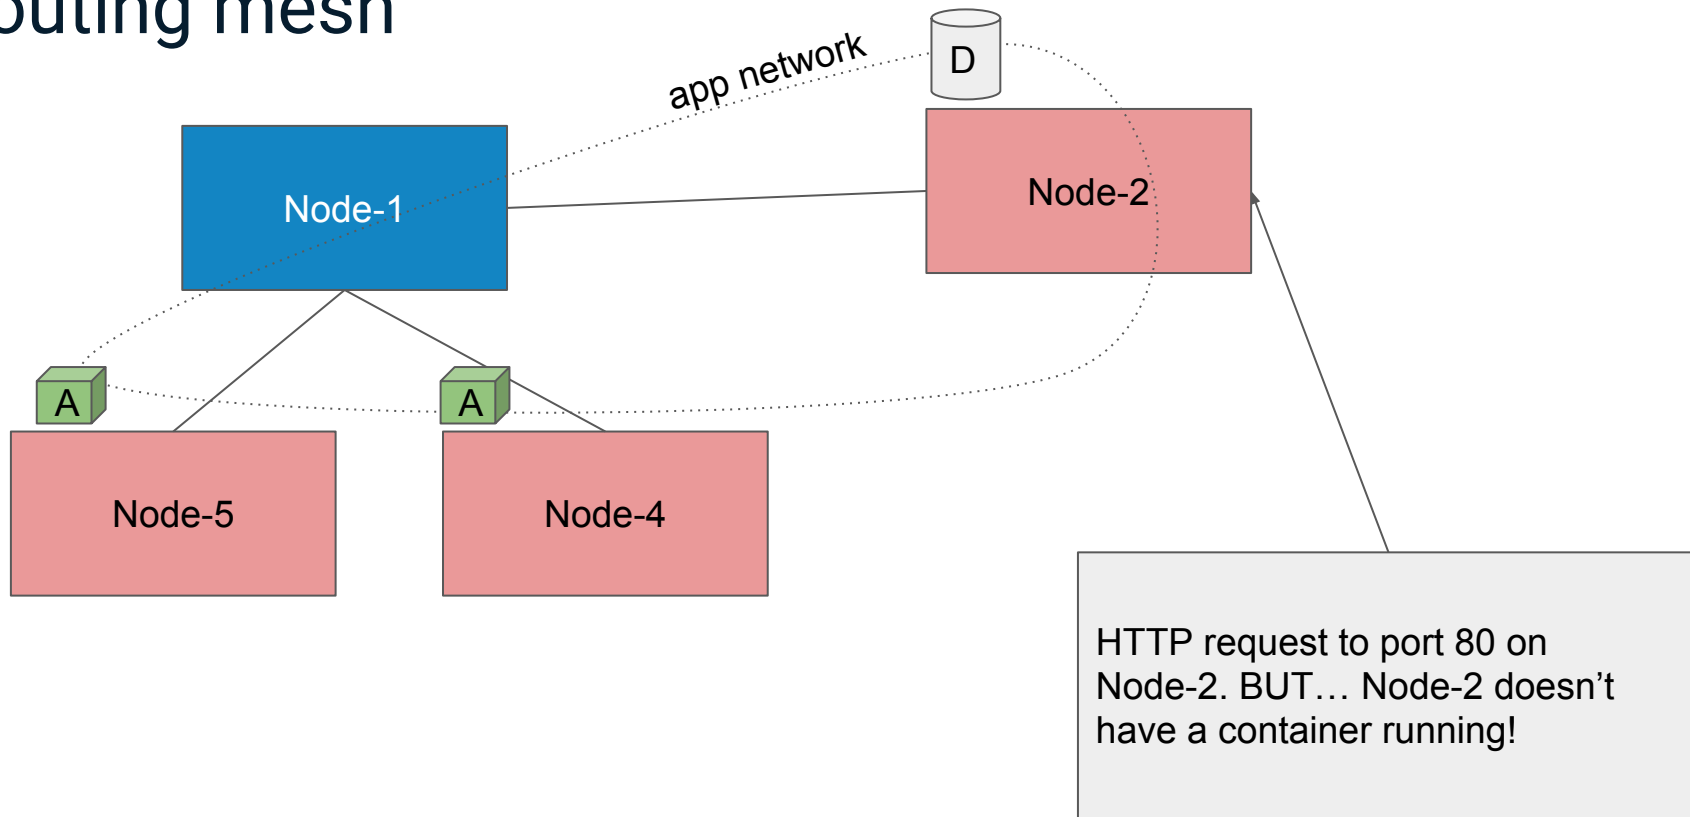


Scale down

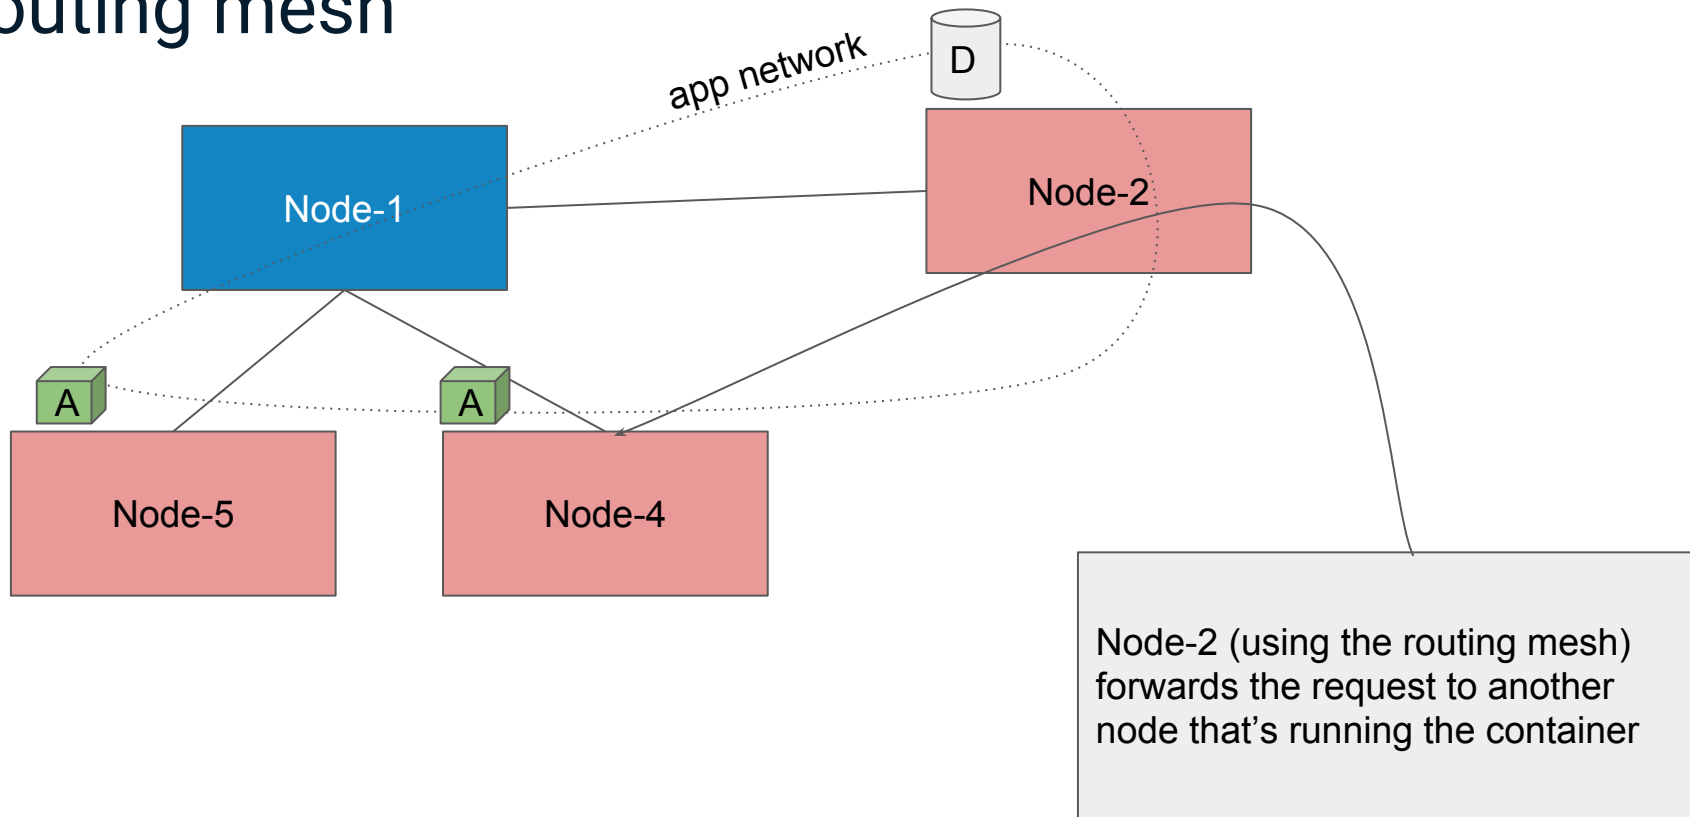


`docker service scale app=2`

Routing mesh

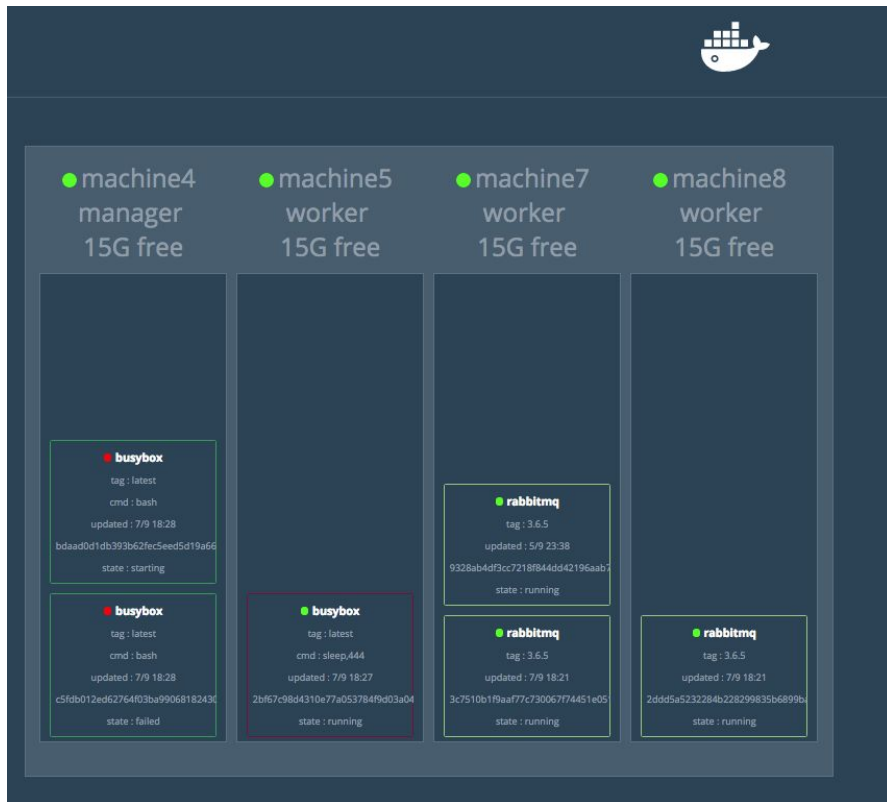


Routing mesh



Swarm Visualizer

- Mano Marks (works at Docker) created a tool to help visualize what's going on in a Swarm
- <https://github.com/ManoMarks/docker-swarm-visualizer>



Swarm recap

- Built-in orchestration that overcomes “fire and forget”
- Requires user to specify a desired state
 - Will use scheduling to get to desired state
 - Self-heals in case of node failure to restore desired state
- Routing mesh ensures traffic gets to a container
 - Allows multiple tasks to run on same machine, even when exposing the same port
- Secure by default

Questions?